

Dish Matrix Recommendations

Answers three questions: why vague Asks look the same today, how much a full Database move helps, and whether to split many sequential filter tables. Opinion is adjusted to how Rasaoi actually works — not a literal copy of the sketch. **Rev 1.2** maps your table-chain **flavor/spice** concern into soft choice dimension **S** inside outcome math J, and tickets it as **ROE-024** (after ROE-023 match rules).

GROUNDING IN
Combined upgrade + registry audit Rev 1.1

Q1
Vague Ask sameness

Q2
Full DB transition

Q3
Tags & gates (not 10 joins)

DOC DATE
7 Sep 2026 · Rev 1.2

STATUS
Recommendation only

Sequence: registry audit Rev 1.1 (\$9b M-01...M-06) → this doc → ticket P0 match rules (**ROE-023**) → then ROE-024+ math/host.

CONTENTS

1. Short answers
2. Q1 — Vague prompts & “same” results
3. Confusion scorecard
- 3b. Named-dish match rules (Rev 1.1)
4. Q2 — Full Database transition
5. Q3 — Sequential tables: opinion
6. Recommended model (hard gates + soft tags)
7. Sketch → Rasaoi mapping
8. Phased advice

01 Short answers

QUESTION	PLAIN ANSWER
1. Vague prompt → same outcomes?	HIGH CONFUSION RISK when the Ask is mood-only and the local menus are thin or similar. Not “AI broken” — soft dials + weak catalog signal.
2. Go fully to Database?	LARGE WIN for freshness, payload, and phone speed. MEDIUM ONLY for vague-Ask variety unless you also add tags + score math. DOES NOT fix bad named-dish match rules by itself.
3. Many sequential tables?	DO NOT build a mandatory join chain (location→gender→health→...→dish). Use hard gates + soft tags on live dishes instead.
4. Named dish → correct hero?	NOT GUARANTEED TODAY — soft token F() can crown Butter Dosai for “Butter chicken.” Fix match rules (P0) before celebrating math moves.

Keep forever: real menus only, catalog plate gate, sanitize after AI, Ask-fulfillment first *with honest F()*, Lab promote gate, no clinical over-claims.

02 Q1 — Vague prompts & “same” results

When someone says “something nice” or only moves mood dials, the system has little hard signal. Dials sit near the middle. Many EDH/Folsom kitchens share similar North Indian staples. Ranking then leans on soft vibe math — so Readings look alike.

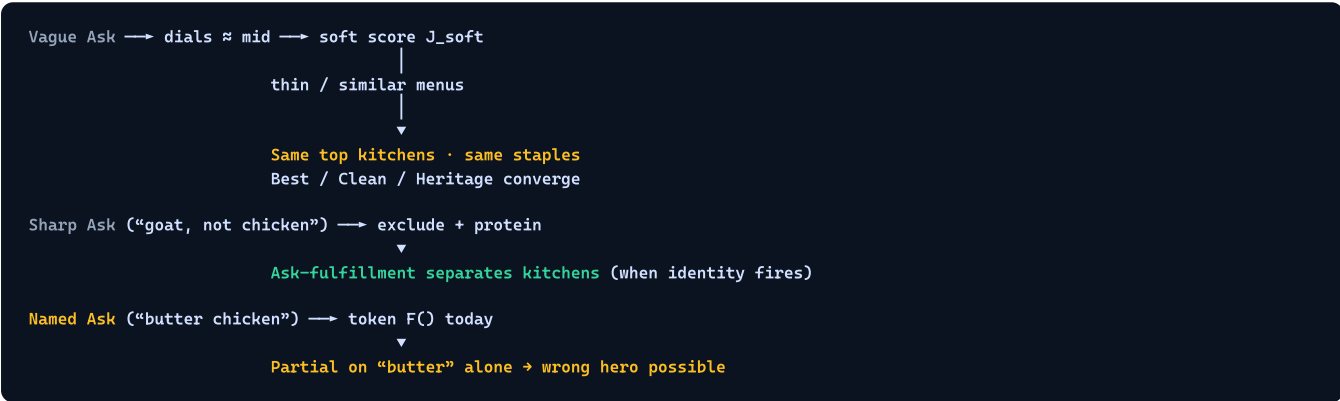
Why it feels the same (today)

ASK SIDE

Vague text → default dials. Few filters. No named dish, protein, or exclusion. Soft scores dominate.

CATALOG SIDE

Overlapping menus. Identity enrich improved (proteins/diet ~100% post ROE-022); cuisine_region ~30%, gi_band ~0% remain. **Match-rule holes** still allow soft F() false positives.

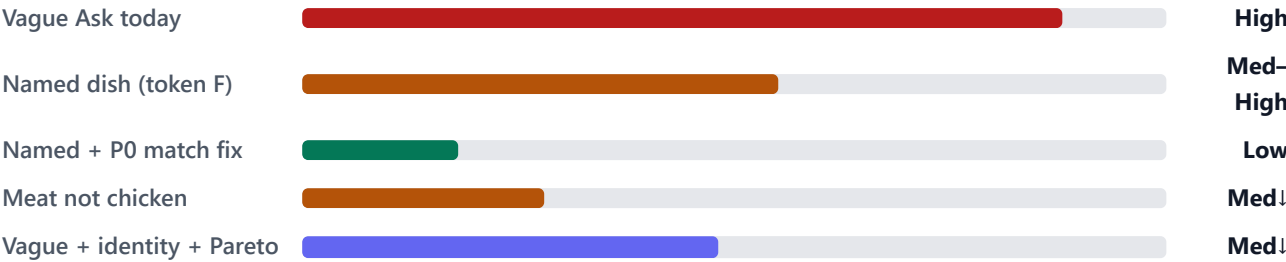


Important: “Results say the same” is often *catalog poverty* + *soft ranking*, not a broken model. Separately, *named-dish false heroes* are **match-rule debt** (\$3b) — not fixed by inventing plates or by J weights alone.

What would make Readings feel different

CHANGE	HOW IT HELPS
P0 tight named-dish / F() rules	Butter chicken → exact kitchen; honesty when only garnish overlap
Richer dish identity (proteins, food_type, region)	Even soft Asks can tilt by region / lightness without inventing plates
Pareto triple plates (upgrade)	Best / Clean / Heritage optimize different axes — less clone slots <i>after</i> F is honest
Softmax + cuisine diversity for Alternates	Hero + alts stop looking like twins
Live DB menus (fresher, more items)	More real choices to pick among
Honest empty / Closest path	Better than fake 100% on weak partials

Illustrative “sameness” / confusion of top-3 (lower = more distinct / correct)



Rev 1.1: named-dish bar raised after Butter chicken → Butter Dosai soak · not a live A/B test.

03 Confusion scorecard (by Ask type)

ASK TYPE	EXAMPLE	CONFUSION TODAY	WHY	FIX LEVERAGE
Vibe-only	“Something nice / comfort food”	HIGH	Default dials; soft score; similar staples	Pareto + diversity + better tags
Wellness vague	“Something healthy”	MED-HIGH	Purity up; no fake “Healthy” cuisine; menus still heavy	dish_type / purity tags on dishes
Protein + exclude	“Meat, not murgi”	MED↓	Identity proteins ~100% post ROE-022; still watch vernacular aliases	Maintain enrich CI + exclusion soak
Named dish	“Butter chicken” / “Mysore pak”	MED-HIGH (Rev 1.1)	Was documented as Low — wrong. Soft token F() + honesty only on none ; soak: Butter Dosai @ 100%	P0 match rules (§3b); then fresher menus
Ethic / diet hard	“Jain” / “Halal”	LOW if tagged	Hard gate drops non-compliant	Diet/ethic tags on dishes + kitchens

BEFORE (VAGUE)

Three kitchens ~same %. Best often dal/naan-adjacent. Clean ≈ lighter fry. Heritage ≈ another curry. Feels cloned.

AFTER (IDENTITY + HONEST F + PARETO)

Same real menus, but F is honest on named Asks; slots pull apart: Ask-align Best, lighter Clean, regional Heritage. Alternates diversify cuisine.

03b Named-dish match rules (Rev 1.1)

Registry audit §9b lists holes M-01 ... M-06 . This section turns them into product rules. Soft score J must *not* override a failed named-dish gate.

Soak: Butter chicken under 40 → Mylapore / Butter Dosai 100% while India Oven had Butter Chicken. Catalog was fine; **F() token policy** was not.

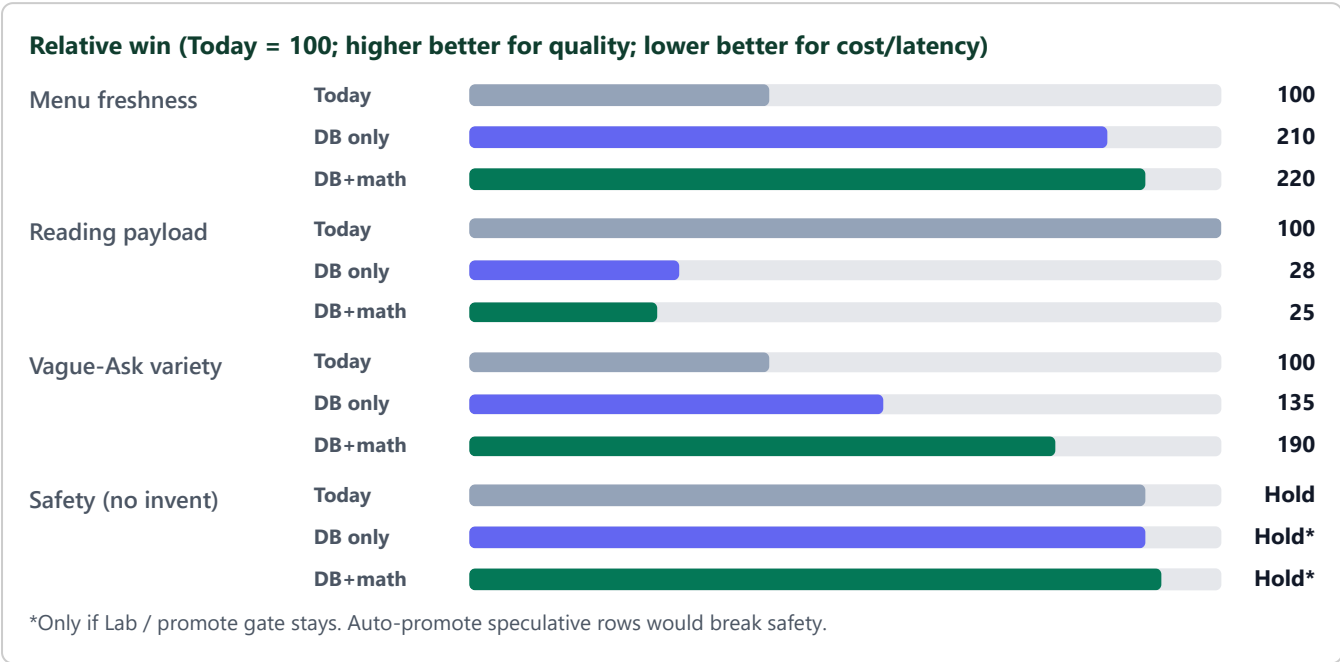
RULE	BEHAVIOR	MAPS TO
R1 Protein required	If Ask tokens include a protein (chicken, goat, paneer...), exact/strong F requires that protein (or synonym) on the plate line	M-01, M-05
R2 Fat/garnish alone ≠ dish	butter / ghee overlap without other Ask tokens → Closest / partial at most; never naked ~100%	M-01, M-02, M-06
R3 Honesty on weak partial	Extend ROE-018: cap confidence + Closest copy when match is garnish-only, not only when none	M-02
R4 No full F from weak doubles	venueAskFulfillment must not promote “full” from two weak partials	M-03
R5 Exact kitchen first	Any kitchen with exact named dish outranks garnish-only kitchens before purity / Sovereign soft tags	M-04
R6 Soak pack	Vitest: Butter chicken; Butter chicken under 40; ghee roast (legitimate); Mysore pak (no dosa false hit)	Acceptance

Ticket: ROE-023 implements R1–R6 in dishIntent.ts / askFulfillment.ts / veda.ts + tests. Do *not* wait for score-reading to “fix” this.

04 Q2 — Full Database transition

“Full DB” here means: **live culinary facts** + `menu_items` **feed scoring** (combined proposal: live cache + `score-reading`). It does *not* mean delete the matrix overnight or auto-promote Apify scrapes.

Today (SPA + static index) DB catalog only DB + upgrade math (J / Pareto / tags)



WIN	DB CATALOG ONLY	DB + UPGRADE MATH	MAGNITUDE
Freshness / fewer honest misses	Yes	Yes + identity for F()	LARGE
Payload / phone CPU	Yes (<code>score-reading</code>)	Same	LARGE
Vague-Ask diversity	Some (more real dishes)	Much better (Pareto + tags)	MED → LARGE
Explain “why this plate”	Weak	J component breakdown	LARGE WITH MATH
Safety	Same if Lab gate	Same	HOLD

Caveat: Database alone does *not* fix vague-Ask sameness, and does *not* fix named-dish token false heroes. You still need honest F() match rules, dish attributes (tags/identity), hard gates, and ranking math (J / Pareto).

05 Q3 — Sequential tables: opinion

The sketch (Ask → location → event → gender → user type → health → ethic → diet → macro → therapeutic → dish) is a useful *thinking list*. As a literal SQL join pipeline, it is a poor fit for Rasaoi.

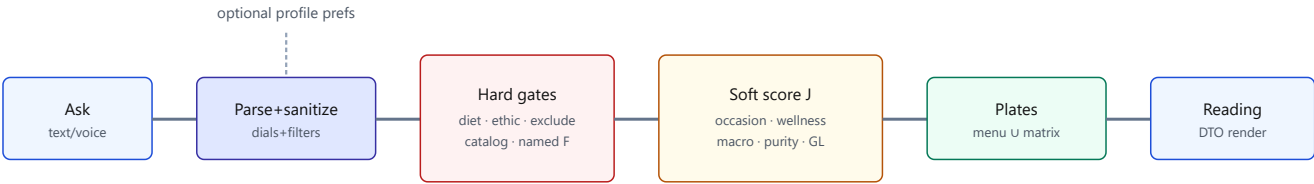
Why not a 10-step join chain

PROBLEM	WHAT GOES WRONG
Most Asks omit stages	Gender, age, renal, DASH rarely said. Empty stages = noise and ops cost.
Hard AND filters shrink Eligible	Too many gates → empty set → same fallbacks / “no match” loops (more sameness).
Gender / age as dish filters	EDH menus are not tagged this way; product and fairness risk if forced as hard gates.
Clinical health tables	“Diseased / renal / pregnant” as ranked DB paths invites medical claims. Handbook: wellness, not clinical.
One mega-column table	Brittle migrations; every new diet flag becomes a column war.

Recommendation: treat the sketch as a **dimension checklist**, not as a mandatory pipeline of tables in the hot path.

06 Recommended model (hard gates + soft tags)

Keep the Outcome Engine path simple. Split *must-pass* from *nice-to-rank*.



ELIGIBLE THEN RANK (SAME IDEA AS UPGRADE J)

Eligible = dishes/kitchens that pass diet n ethic n exclude n regional n catalog n **named-Ask F policy**

Maximize $J = w_F \cdot F + w_D \cdot D + w_P \cdot P + w_B \cdot B + w_W \cdot W + w_S \cdot S - w_G \cdot G$

Soft choice dimensions from the sketch (occasion, macro, **flavor/spice**, wellness, household) feed $W / S / D$ – they do not each require a join table, and they do not override hard gates or named F (Rev 1.2).

Data shape (lean)

```
dishes / culinary knowledge (live rows)
├── dish_tags (dish_id, dimension, tag, speculation_tier)
│   └── dimension ∈ { diet, ethic, macro, wellness, spice, therapeutic, allergen, food_type, ... }
└── Optional profile (prefs only):
    ├── location · household (kid/adult) · default occasion
    └── Bio / clinical raw signals stay on-device / caveated

Not recommended: gender_only table as hard dish filter;
                  "diseased" as a ranked therapeutic table without legal copy
```

LAYER	EXAMPLES	BEHAVIOR
Hard gates	Jain, vegan, halal, allergy excludes, catalog membership, named-Ask protein/F policy (R1–R5)	Fail → out or Closest-capped. No soft override to naked 100%.
Soft score	Celebration dial, light/fresh tags, keto-ish, purity, budget, GL lens, mild/non-spicy (S)	Reorder Eligible; never invent dishes.
Profile prefs	EDH/Folsom, family night, prefer veg	Defaults for parse — user Ask still wins.

07 Sketch → Rasaoi mapping

Your concepts, adjusted to what already exists (dials, DIET-001, exclude_ingredients, identity, GL lens).

YOUR SKETCH	MAPS TO TODAY	HARD OR SOFT?	RECOMMENDED STORAGE
Ask → intent parsing	<code>parse-intent</code> + sanitize	—	Keep; always sanitize after LLM
Location table	EDH/Folsom venues; Places	Soft / scope	Restaurant geo + profile default city
Event / occasion	Context dial + <code>context_tags</code>	Soft	Dial + optional dish tags (celebratory...)
Gender table	— (not in product)	AVOID HARD	Do not gate dishes by gender
User type (kid/adult/elderly)	Twin / prefs (thin today)	Soft pref	Profile pref; mild score tilt only if tagged
Health (allergy...)	<code>exclude_ingredients</code> ; GL lens	Allergy = hard; GL = soft	Allergen tags + excludes; lens caveats
Health (pregnant / diseased...)	Not clinical today	DEFER	Only with legal copy + verified tags; never invent care advice
Ethic (halal, jain...)	Strict dietary / certifications	Hard	Kitchen + dish tags (DIET-001 family)
Diet (vegan, veg, non-veg...)	<code>diet_class</code> + dietary gates	Hard	On <code>dishes</code> / identity; junction tags OK
Macro (keto, high-protein...)	Wellness / purity / macros	Soft	Tags + nutrition fields; <code>speculation_tier</code>
Flavor / spice (mild, non-spicy, hot...)	Stop-words today — not scored (soak: “chicken, non spicy”)	Soft (Rev 1.2)	<code>spice_level</code> / <code>dish_tags.spice</code> + Ask pref in J (w_S-S); name heuristics until tags dense
Therapeutic (DASH, Low-FODMAP...)	Experimental / sparse	Soft + caveated	Optional tags after verified enrich; wellness copy only
Dish table	<code>menu_items</code> + <code>dishes</code> + index	Source of plates	Live DB primary; matrix/index enrich until cache ships

08 Phased advice — Rev 1.2

- 1. **P0 — Named-dish / F() match rules — ROE-023** (R1–R6). *Blocks false 100% heroes.*
- 2. **Choice dimensions in J (spice/flavor first) — ROE-024**: stop ignoring mild/non-spicy; soft *S* in outcome math; tag/heuristic *spice_level*; keep protein Ask-align on all triple slots.
- 3. **Maintain identity + alias parity** — *cuisine_region / food_type*; *culinary:check-aliases*.
- 4. **Shared scoring + live cache + score-reading — ROE-025+**: host the *same* J (incl. *S*) on Edge; Lab gate.
- 5. **Full tag model** — *dish_tags* *ethic / diet / macro / wellness / spice / therapeutic* + *speculation_tier*.
- 6. **Vague-Ask UX** — Pareto + “why” bars; honest Closest.
- 7. **Defer** — gender hard-gates; clinical therapeutic ranking.

Do now ROE-023 match · ROE-024 spice/S in J	Do next Shared J on Edge · live cache	Do with care Full tags · GL · therapeutic	Don't 10-table AND · invent plates · math without S/F
--	---	---	--

Bottom line: fix sameness with *better real choices + clearer gates + soft ranking math*. Fix named-dish false heroes with *honest F()*. Fix “non spicy” with *choice-dimension S in J* — not a 10-table join. Move cognition to the Database for speed — encode the sketch as tags and weights.